

# Model Stealing Attack

Byung Jae Bae

December 17, 2023

## Abstract

We perform a model stealing attack on a pretrained Resnet50 model by using the KnockoffNet methodology.

## 1 Introduction

We run a KnockoffNet attack on a pretrained Resnet50 model trained on the CIFAR dataset with another Resnet50 model but trained on both CIFAR and MNIST datasets. We test the new KnockoffNet trained on 10k, 20k, 30k, 40k, and 50k queries to the original pretrained model on the CIFAR test dataset, MNIST test dataset, and both combined CIFAR and MNIST test dataset to compare the how the different datasets affect the test accuracies.

## 2 Methods

We used a pretrained Resnet50 model trained on CIFAR dataset from hugging-face as our victim model. We created another Resnet50 model with randomized weights to perform a model stealing attack using the KnockoffNet method. The training data was a mix of MNIST and CIFAR data. We recorded the accuracies of

## 3 Results

The accuracies of the CIFAR test dataset for 10k, 20k, 30k, 40k, and 50k queries were 0.137, 0.1154, 0.151, 0.1797, and 0.161, respectively. For the MNIST and CIFAR dataset it was 0.1984, 0.18735, 0.18895, 0.23545, and 0.21055. For only MNIST it was 0.189, 0.1795, 0.1488, 0.2127, 0.2185. The table below makes a clearer illustration.

## 4 Discussion

The results were somewhat surprising. Firstly, the accuracies for all test accuracies were very low even for 50k queries. Secondly, there were no linear trends

Number of Queries	CIFAR	MNIST and CIFAR	MNIST
10000	0.137	0.1984	0.189
20000	0.1154	0.18735	0.1795
30000	0.151	0.18895	0.1488
40000	0.1797	0.23545	0.2127
50000	0.161	0.21055	0.2185

Figure 1: Table of accuracies for different number of queries.




Figure 2: Graph of accuracies for different number of queries.

between number of queries and test accuracies, as the test accuracies would fluctuate throughout the different query numbers. Despite no linear trends the overall trend seems still to trend positively which would be consistent with our expectations. One aspect of our training that was imposed was to use MNIST and CIFAR to replicate a model only trained on CIFAR. This most likely have affected our testing accuracies because our Knockoff Model was training on out of distribution data compared to what the original model was made for. Since we first feed the MNIST data to the pretrained model only trained on CIFAR, the output would likely be incorrect as it was not trained for the task, but downstream our model was fed the incorrect output leading to a worse model to generalize to MNIST. This is why MNIST almost consistently scored the lowest accuracy out of the 3 test datasets.